

# CIFAR-10 image recognition task

# The CIFAR-10 data

## Data splits

The data (n=60,000) was split into a training set (n=50,000) and a validation set (n=10,000)

## Rationale for validation

Neural networks are prone to over-fit the training data (memorise)

Performance of the model must be assessed in validation data that the network has not seen during training

## Data structure

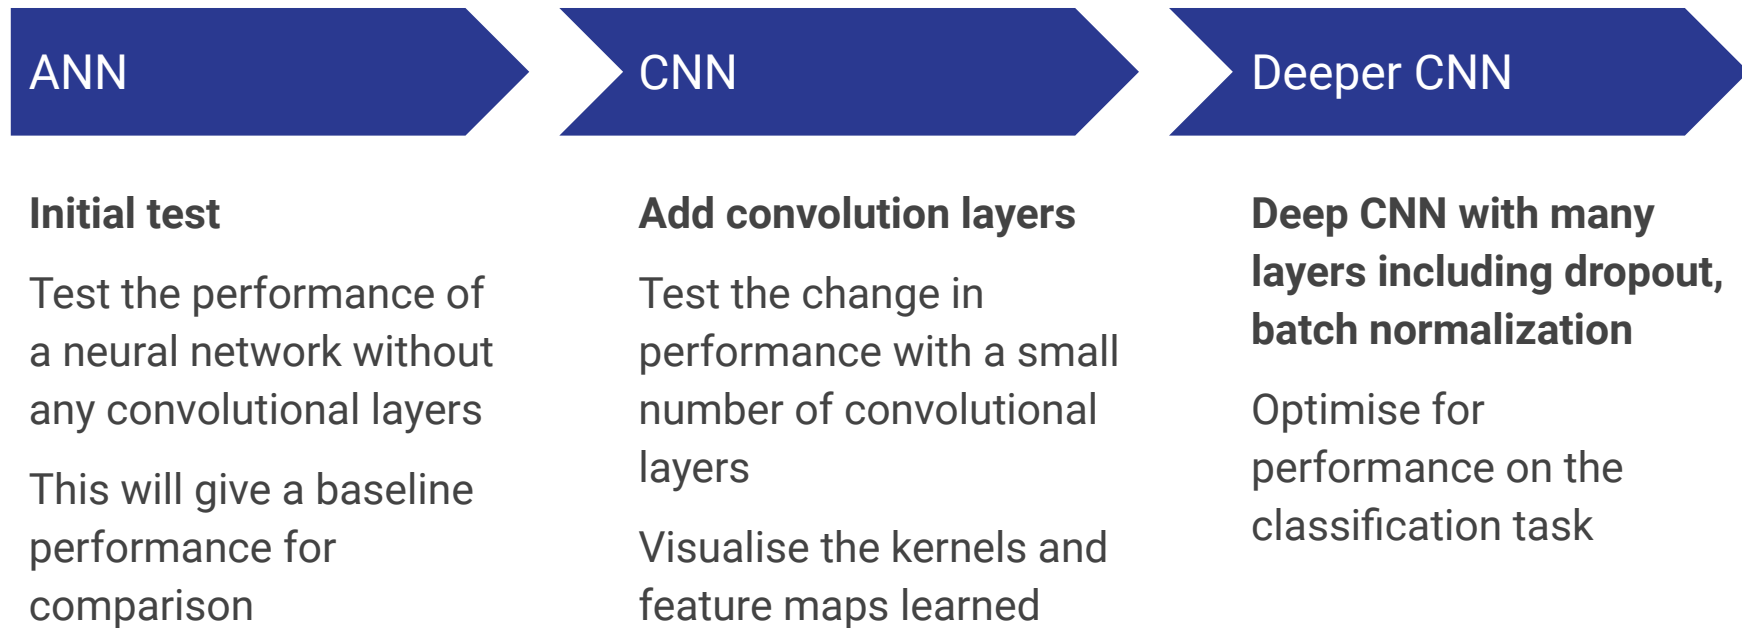
Shape of each data instance was (32, 32, 3)

32 x 32 pixels, three channels (RGB)

The values were normalised by division by 255 (to range 0-1)

10 classes (labels)

# Approach to model building experiments



# Modelling choices

## Loss function

### **Sparse categorical cross entropy loss**

This is appropriate for a multi-class classification objective

## Metrics

### **Accuracy**

This is appropriate for a balanced multi-class problem

**NB** if classes were imbalance this would not be appropriate

## Optimizer

### **Stochastic gradient descent**


Ten epochs

Chosen empirically as in initial models, accuracy plateaus before ten epochs

# ANN: from one to three FC layers


Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 10)	30730

=====  
Total params: 30,730  
Trainable params: 30,730  
Non-trainable params: 0  
=====



Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 32, 32, 3)]	0
flatten_2 (Flatten)	(None, 3072)	0
dense_3 (Dense)	(None, 100)	307300
dense_4 (Dense)	(None, 50)	5050
dense_5 (Dense)	(None, 10)	510

=====  
Total params: 312,860  
Trainable params: 312,860  
Non-trainable params: 0  
=====




# CNN: one convolutional layer




- 1 convolutional layer
- 50 kernels
- Kernel size = (3,3)
- Stride = 1
- Activation = ReLU
- Optimizer = Adam (Kingma et al. 2015)
- Metric = accuracy
- 1 FC layer for classification

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d (Conv2D)	(None, 32, 32, 50)	1400
flatten_3 (Flatten)	(None, 51200)	0
dense_6 (Dense)	(None, 10)	512010

=====  
Total params: 513,410  
Trainable params: 513,410  
Non-trainable params: 0  
=====



# CNN: one convolutional layer



cat

# CNN: 3 convolutional layers

- 3 convolutional layers

- 50, 100, 200 kernels

- Kernel (3,3) stride 1

- Max pooling layer size (2,2)


- 3 FC layers: 40, 40, 10

- Activation = ReLU

- Optimizer = Adam




- Metric = accuracy

Layer (type)	Output Shape	Param #
input_19 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d_39 (Conv2D)	(None, 32, 32, 50)	1400
max_pooling2d_24 (MaxPooling2D)	(None, 16, 16, 50)	0
conv2d_41 (Conv2D)	(None, 16, 16, 100)	45100
max_pooling2d_25 (MaxPooling2D)	(None, 8, 8, 100)	0
conv2d_42 (Conv2D)	(None, 8, 8, 200)	180200
max_pooling2d_26 (MaxPooling2D)	(None, 4, 4, 200)	0
flatten_14 (Flatten)	(None, 3200)	0
dense_32 (Dense)	(None, 40)	128040
dense_33 (Dense)	(None, 40)	1640
dense_34 (Dense)	(None, 10)	410
=====		
Total params: 356,790		
Trainable params: 356,790		
Non-trainable params: 0		







# CNN: 3 convolutional layers




1st convolutional layer

Max pooling



2nd convolutional layer

Max pooling



3rd convolutional layer


Max pooling



# CNN: adding dropout

Layer (type)	Output Shape	Param #
input_37 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d_144 (Conv2D)	(None, 32, 32, 50)	1400
max_pooling2d_88 (MaxPooling2D)	(None, 16, 16, 50)	0
conv2d_145 (Conv2D)	(None, 16, 16, 100)	45100
max_pooling2d_89 (MaxPooling2D)	(None, 8, 8, 100)	0
conv2d_146 (Conv2D)	(None, 8, 8, 200)	180200
conv2d_147 (Conv2D)	(None, 8, 8, 200)	360200
conv2d_148 (Conv2D)	(None, 8, 8, 100)	180100
max_pooling2d_90 (MaxPooling2D)	(None, 4, 4, 100)	0
flatten_36 (Flatten)	(None, 1600)	0
dropout_1 (Dropout)	(None, 1600)	0
dense_98 (Dense)	(None, 40)	64040
dense_99 (Dense)	(None, 40)	1640
dense_100 (Dense)	(None, 10)	410

=====  
Total params: 833,090  
Trainable params: 833,090  
Non-trainable params: 0



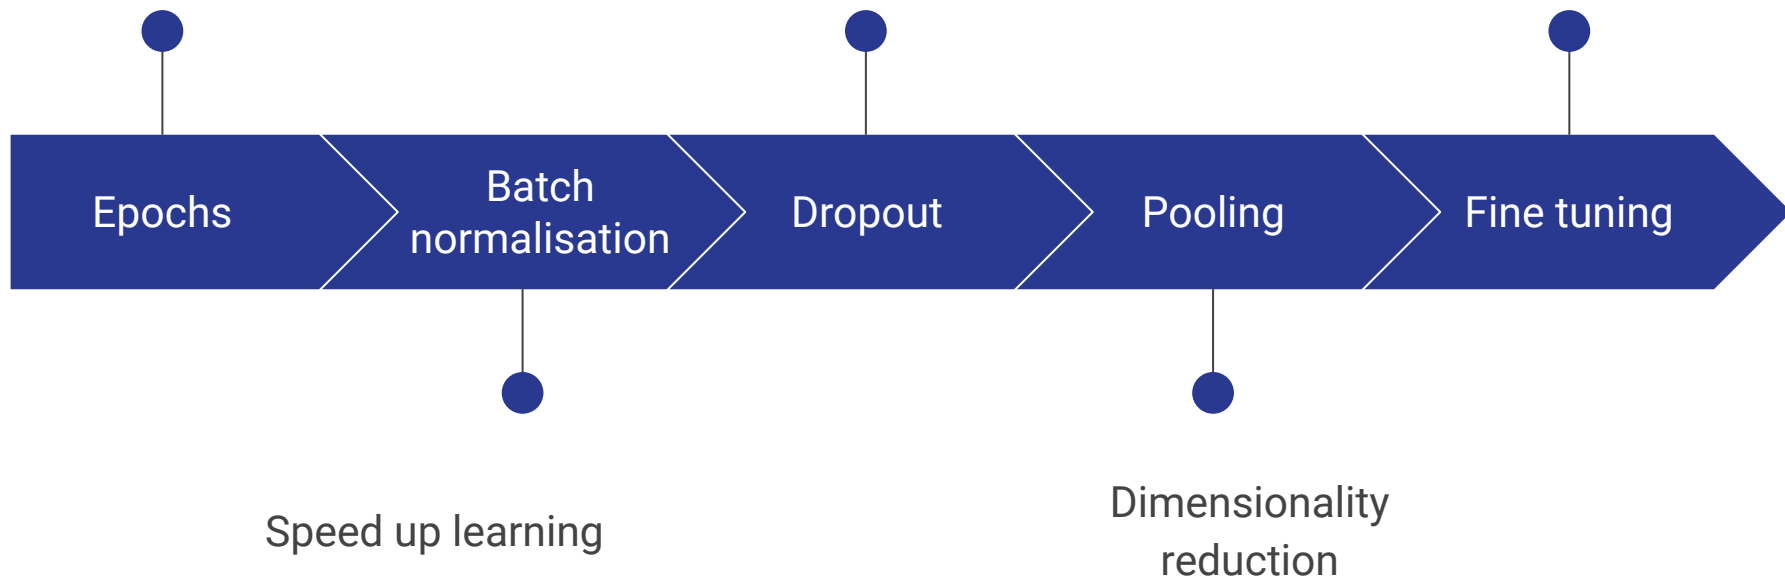
# CNN PERFORMANCE OPTIMISATION

# Optimising for performance

Allowing more time to learn

Simplifying learning


Playing with functions, learning rates etc. and iterating













Model: "sequential"


Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496

# Performance metrics





# Confusion matrix

Airplane	0	
Automobile	1	
Bird	2	
Cat	3	
Deer	4	
Dog	5	
Frog	6	
Horse	7	
Ship	8	
Truck	9	



# Learning curve





# Conclusions

# Conclusions

- CNN performs better than ANN for image data
- Max pooling layers reduce the growth in parameter numbers when adding more layers to CNN
- Dropout layer reduced the over-fitting on training data
- Batch normalisation standardises input and stabilises learning process
- Visualisation of the feature maps of the CNN demonstrates deeper layers learn higher order features and more feature maps show no activation for a given image instance
- We were able to get relatively good training results with a very small number of epochs, but to prevent over-fit and to get to a truly high accuracy, we had to not only apply additional techniques and layers, but also increase significantly number of epochs
- Creating a well performing model requires a lot of skill as well as time and resources for the model to be trained



# References

# References

- Madhugiri, D. (2021) *Using CNN for image classification on CIFAR-10 Dataset*, Medium. Available at: <https://devashree-madhugiri.medium.com/using-cnn-for-image-classification-on-cifar-10-dataset-7803d9f3b983> (Accessed: 22 June 2023).
- Garg, A. (2022) *How to classify the images of the CIFAR-10 dataset using CNN?*, Analytics Vidhya. Available at: <https://www.analyticsvidhya.com/blog/2022/09/how-to-classify-the-images-of-the-cifar-10-dataset-using-cnn/> (Accessed: 22 June 2023).
- Le, J. (2018) *The 4 convolutional neural network models that can classify your fashion images*, Medium. Available at: <https://towardsdatascience.com/the-4-convolutional-neural-network-models-that-can-classify-your-fashion-images-9fe7f3e5399d> (Accessed: 22 June 2023).
- Kubat, M. (2021) *Introduction to machine learning*. Cham: Springer.
- Chollet, F. (2021) 'Getting started with neural networks', in *Deep learning with Python*. Shelter Island: Manning Publications, pp. 56–92.
- *Implementing a CNN in Tensorflow & Keras* (2023) LearnOpenCV. Available at: <https://learnopencv.com/implementing-cnn-tensorflow-keras/> (Accessed: 11 July 2023).
- Chollet, F., 2021. *Deep learning with Python*. Simon and Schuster.
- Kingma, D.P. and Ba, J., 2014. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.

